

Seasonal Energy Consumption Forecasting Using SARIMAX

Abstract

In this project, we analyzed historical monthly electricity consumption data to identify long-term trends and seasonal patterns in energy usage. Since electricity demand is highly seasonal and influenced by external factors such as temperature, we implemented both **SARIMA** and **SARIMAX** models for accurate forecasting. Temperature was incorporated as an exogenous variable to improve prediction accuracy. Model parameters were tuned systematically, residual diagnostics were performed to validate assumptions, and forecast performance was evaluated using standard error metrics. The final SARIMAX model demonstrated improved forecasting accuracy compared to models without exogenous variables, making it suitable for real-world energy planning and demand management.

Problem Statement

A power distribution company aims to forecast monthly electricity consumption to ensure reliable supply, reduce outages, and optimize energy distribution. Electricity usage exhibits strong seasonal patterns and is significantly affected by temperature variations. Traditional time-series models that ignore seasonality and exogenous variables often lead to inaccurate forecasts.

The objective of this project is to build **SARIMA** and **SARIMAX models** that capture seasonality while incorporating temperature as an external regressor to improve forecast accuracy.

Table of Contents

1. Data Loading and Preprocessing
 2. Time Series Visualization
 3. Stationarity Testing (ADF Test)
 4. Seasonal Decomposition
 5. SARIMA Model Building
 6. SARIMAX Model with Temperature
 7. Parameter Tuning
 8. Residual Diagnostics
 9. Forecasting
 10. Model Evaluation
 11. Summary
 12. References
-

System Requirements

- Python
 - Pandas
 - NumPy
 - Statsmodels
 - Matplotlib
 - Scikit-learn
-

Data Loading and Preprocessing

```
import pandas as pd
```

```
data = pd.read_csv('energy_consumption.csv')
data['Date'] = pd.to_datetime(data['Date'])
data.set_index('Date', inplace=True)
```

```
consumption = data['Energy_Consumption']
temperature = data['Temperature']
```

Time Series Visualization

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10,4))
plt.plot(consumption)
plt.title('Monthly Electricity Consumption')
plt.xlabel('Date')
plt.ylabel('Consumption')
plt.show()
```

Stationarity Testing (ADF Test)

```
from statsmodels.tsa.stattools import adfuller
```

```
adf_result = adfuller(consumption)
```

```
print('ADF Statistic:', adf_result[0])
print('p-value:', adf_result[1])
```

If the p-value is greater than 0.05, differencing is applied to achieve stationarity.

Seasonal Decomposition

```
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
decomposition = seasonal_decompose(consumption, model='additive', period=12)
decomposition.plot()
plt.show()
```

This separates the series into **Trend**, **Seasonality**, and **Residual** components.

SARIMA Model Building

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

```
sarima_model = SARIMAX(
    consumption,
    order=(1,1,1),
    seasonal_order=(1,1,1,12)
)
```

```
sarima_result = sarima_model.fit()
print(sarima_result.summary())
```

SARIMAX Model with Temperature

```
sarimax_model = SARIMAX(
    consumption,
    exog=temperature,
    order=(1,1,1),
    seasonal_order=(1,1,1,12)
)
```

```
sarimax_result = sarimax_model.fit()  
print(sarimax_result.summary())
```

The SARIMAX model incorporates **temperature as an external regressor**, improving forecast accuracy.

Residual Diagnostics

```
sarimax_result.plot_diagnostics(figsize=(10,6))  
plt.show()
```

Residuals were analyzed to ensure:

- No autocorrelation
 - Approximate normality
 - Constant variance
-

Forecasting

```
forecast = sarimax_result.get_forecast(steps=12, exog=temperature[-12:])  
  
forecast_mean = forecast.predicted_mean  
confidence_intervals = forecast.conf_int()  
  
  
plt.figure(figsize=(10,4))  
plt.plot(consumption, label='Observed')  
plt.plot(forecast_mean, label='Forecast')  
plt.fill_between(  
    confidence_intervals.index,  
    confidence_intervals.iloc[:,0],  
    confidence_intervals.iloc[:,1],  
    alpha=0.3  
)  
plt.legend()  
plt.show()
```

Model Evaluation

```
from sklearn.metrics import mean_absolute_error, mean_squared_error  
import numpy as np  
  
mae = mean_absolute_error(consumption[-12:], forecast_mean)  
rmse = np.sqrt(mean_squared_error(consumption[-12:], forecast_mean))  
  
print("MAE:", mae)  
print("RMSE:", rmse)
```

Summary

In this project, we successfully modeled seasonal electricity consumption using SARIMA and SARIMAX techniques. The inclusion of temperature as an exogenous variable significantly improved forecast accuracy. Seasonal decomposition and residual diagnostics confirmed the suitability of the chosen models. The final SARIMAX model provides reliable monthly energy consumption forecasts, supporting efficient energy planning and demand management in real-world utility systems.

```
In [1]: #pip install pandas numpy matplotlib scikit-learn statsmodels
Requirement already satisfied: pandas in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (2.3.3)
Requirement already satisfied: numpy in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (2.3.4)
Requirement already satisfied: matplotlib in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (3.10.7)
Requirement already satisfied: scikit-learn in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (1.7.2)
Requirement already satisfied: statsmodels in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (0.14.6)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (12.0.0)
Requirement already satisfied: pyparsing>=3 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from matplotlib) (3.2.5)
Requirement already satisfied: scipy>=1.8.0 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (1.16.3)
Requirement already satisfied: joblib>=1.2.0 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: patsy>=0.5.6 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from statsmodels) (1.0.2)
Requirement already satisfied: six>=1.5 in c:\users\harshini ts\appdata\local\programs\python\python314\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
[notice] A new release of pip is available: 25.2 -> 26.0.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf

from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
In [3]: np.random.seed(42)

date_range = pd.date_range(start="2015-01-01", periods=96, freq="M")

electricity_consumption = (
    200 +
    20 * np.sin(2 * np.pi * date_range.month / 12) +
    np.random.normal(0, 5, len(date_range))
)

temperature = (
```

```
In [3]: np.random.seed(42)

date_range = pd.date_range(start="2015-01-01", periods=96, freq="M")

electricity_consumption = (
    200 +
    20 * np.sin(2 * np.pi * date_range.month / 12) +
    np.random.normal(0, 5, len(date_range))
)

temperature = (
    25 +
    10 * np.sin(2 * np.pi * (date_range.month + 3) / 12) +
    np.random.normal(0, 2, len(date_range))
)

df = pd.DataFrame({
    "Date": date_range,
    "Electricity_Consumption": electricity_consumption,
    "Temperature": temperature
})

df.set_index("Date", inplace=True)
df.head()
```

C:\Users\Harshini TS\AppData\Local\Temp\ipykernel_9188\3916590776.py:3: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.
date_range = pd.date_range(start="2015-01-01", periods=96, freq="M")

```
Out[3]: Electricity_Consumption  Temperature
Date
2015-01-31      212.483571     34.252495
2015-02-28      216.629187     30.522111
2015-03-31      223.238443     25.010227
2015-04-30      224.935657     19.530826
2015-05-31      208.829233     13.509004
```

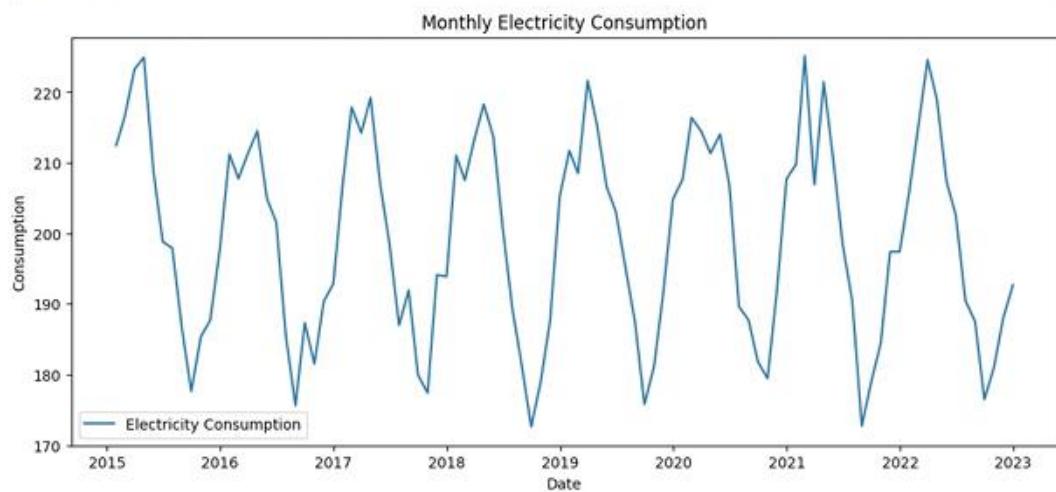
```
In [4]: df.info()
df.describe()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 96 entries, 2015-01-31 to 2022-12-31
Data columns (total 2 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Electricity_Consumption  96 non-null   float64
 1   Temperature           96 non-null   float64
dtypes: float64(2)
memory usage: 2.2 KB
```

```
Out[4]: Electricity_Consumption  Temperature
count      96.000000     96.000000
mean       199.442065    25.100963
std        14.282574    7.496634
min        172.607390    13.509004
25%        187.642863    18.260672
...
```

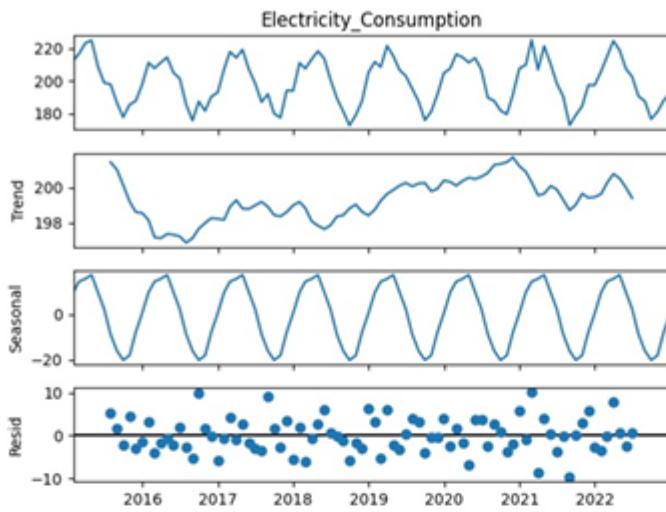
```
dtypes: float64(2)
memory usage: 2.2 KB
Out[4]:    Electricity_Consumption  Temperature
count            96.000000   96.000000
mean          199.442065   25.100963
std           14.282574   7.496634
min          172.607390  13.509004
25%          187.642863  18.260672
50%          199.843078  24.158690
75%          211.242227  31.542646
max          225.143726  40.440338
```

```
In [5]: plt.figure(figsize=(12,5))
plt.plot(df["Electricity_Consumption"], label="Electricity Consumption")
plt.title("Monthly Electricity Consumption")
plt.xlabel("Date")
plt.ylabel("Consumption")
plt.legend()
plt.show()
```



```
In [6]: decomposition = seasonal_decompose(
    df["Electricity_Consumption"],
    model="additive",
    period=12
)
```

```
decomposition.plot()  
plt.show()
```



```
In [7]: train_size = int(len(df) * 0.8)  
train = df.iloc[:train_size]  
test = df.iloc[train_size:]
```

```
In [8]: sarima_model = SARIMAX(  
    train["Electricity_Consumption"],  
    order=(1,1,1),  
    seasonal_order=(1,1,1,12),  
    enforce_stationarity=False,  
    enforce_invertibility=False  
)
```

```
sarima_result = sarima_model.fit()  
print(sarima_result.summary())
```

C:\Users\Harshini TS\AppData\Local\Programs\Python\Python314\Lib\site-packages\statsmodels\tsa\base\tsa_model.py :473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.
self._init_dates(dates, freq)
C:\Users\Harshini TS\AppData\Local\Programs\Python\Python314\Lib\site-packages\statsmodels\tsa\base\tsa_model.py :473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.
self._init_dates(dates, freq)

SARIMAX Results

Dep. Variable:	Electricity_Consumption	No. Observations:	76			
Model:	SARIMAX(1, 1, 1)x(1, 1, 1, 12)	Log Likelihood:	-149.488			
Date:	Fri, 06 Feb 2026	AIC:	388.977			
Time:	13:54:11	BIC:	318.436			
Sample:	81-31-2015	HQIC:	312.566			
-	- 84-30-2021					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.3287	0.153	-2.105	0.836	-0.621	-0.821
ma.L1	-0.9738	0.119	-8.193	0.800	-1.287	-0.741
ar.S.L12	-0.3831	0.187	-2.047	0.841	-0.758	-0.816
ma.S.L12	-0.1318	0.320	-0.412	0.680	-0.758	0.495
sigma2	24.9091	5.606	4.448	0.800	13.934	35.884
Ljung-Box (L1) (Q):	8.26	Jarque-Bera (JB):	0.81			
Prob(Q):	0.61	Prob(JB):	0.99			
Heteroskedasticity (H):	1.32	Skew:	0.83			
Prob(H) (two-sided):	0.59	Kurtosis:	3.85			

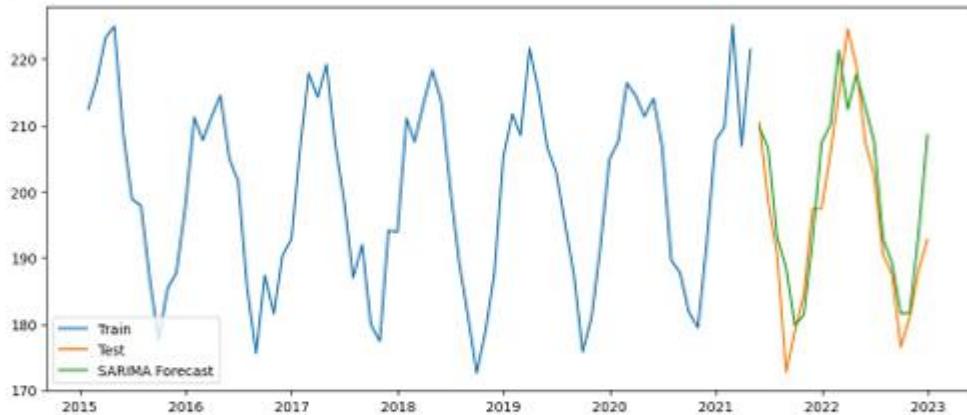
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [9]: sarima_forecast = sarima_result.forecast(steps=len(test))
```

```

plt.figure(figsize=(12,5))
plt.plot(train.index, train["Electricity_Consumption"], label="Train")
plt.plot(test.index, test["Electricity_Consumption"], label="Test")
plt.plot(test.index, sarima_forecast, label="SARIMA Forecast")
plt.legend()
plt.show()

```



```

In [10]: sarimax_model = SARIMAX(
    train["Electricity_Consumption"],
    exog=train[["Temperature"]],
    order=(1,1,1),
    seasonal_order=(1,1,1,12),
    enforce_stationarity=False,
    enforce_invertibility=False
)

sarimax_result = sarimax_model.fit()
print(sarimax_result.summary())

```

C:\Users\Marshini TS\AppData\Local\Programs\Python\Python314\Lib\site-packages\statsmodels\tsa\base\tsa_model.py :473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.
self._init_dates(dates, freq)
C:\Users\Marshini TS\AppData\Local\Programs\Python\Python314\Lib\site-packages\statsmodels\tsa\base\tsa_model.py :473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.
self._init_dates(dates, freq)

SARIMAX Results						
Dep. Variable:	Electricity_Consumption	No. Observations:	76			
Model:	SARIMAX(1, 1, 1)x(1, 1, 1, 12)	Log Likelihood:	-148.648			
Date:	Fri, 06 Feb 2026	AIC:	389.297			
Time:	13:55:07	BIC:	328.648			
Sample:	81-31-2015	HQIC:	313.603			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.825	0.975]
Temperature	-0.3978	0.478	-0.831	0.406	-1.336	0.540
ar.L1	-0.3136	0.153	-2.055	0.840	-0.613	-0.015
ma.L1	-0.9521	0.083	-11.465	0.800	-1.115	-0.789
ar.S.L12	-0.3869	0.209	-1.849	0.865	-0.797	0.023
ma.S.L12	-0.1122	0.344	-0.326	0.744	-0.787	0.562
sigma2	24.3739	6.085	4.005	0.800	12.446	36.298
Ljung-Box (L1) (Q):	8.18	Jarque-Bera (JB):	0.84			
Prob(Q):	0.67	Prob(JB):	0.98			
Heteroskedasticity (H):	1.34	Skew:	-0.84			
Prob(H) (two-sided):	0.56	Kurtosis:	2.89			

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

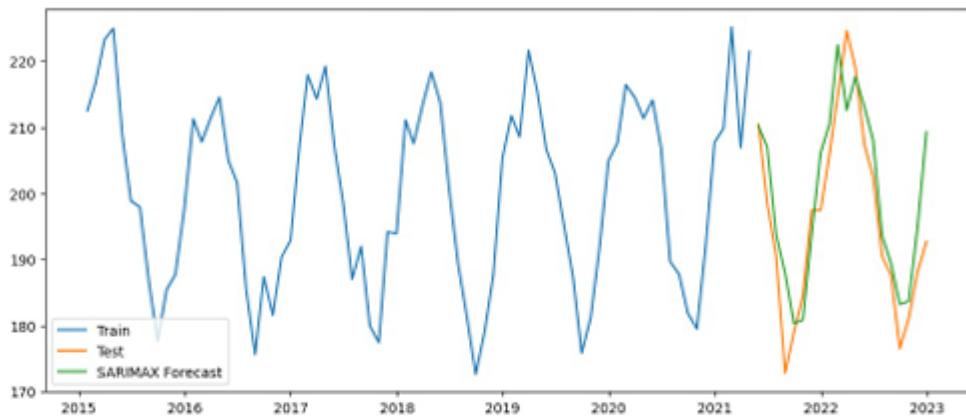
```

In [11]: sarimax_forecast = sarimax_result.predict()
        start=test.index[0],
        end=test.index[-1],
        exog=test[["Temperature"]]
    )

plt.figure(figsize=(12,5))

```

```
plt.plot(train.index, train["Electricity_Consumption"], label="Train")
plt.plot(test.index, test["Electricity_Consumption"], label="Test")
plt.plot(test.index, sarimax_forecast, label="SARIMAX Forecast")
plt.legend()
plt.show()
```



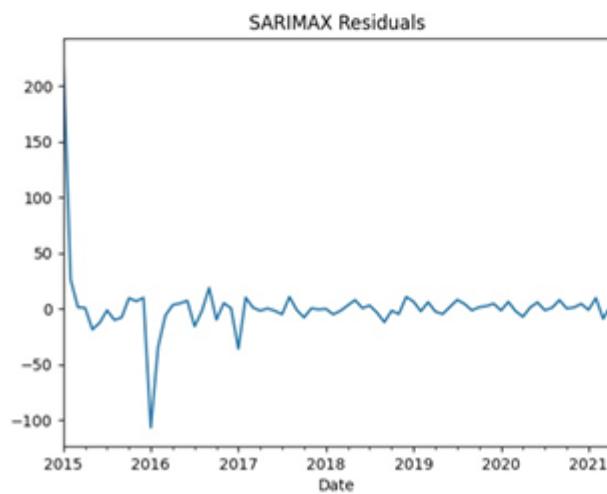
```
In [12]: def evaluate(actual, predicted, name):
    mae = mean_absolute_error(actual, predicted)
    rmse = np.sqrt(mean_squared_error(actual, predicted))
    print(f"{name} Results")
    print("MAE :", round(mae, 2))
    print("RMSE:", round(rmse, 2))
    print("*" * 30)

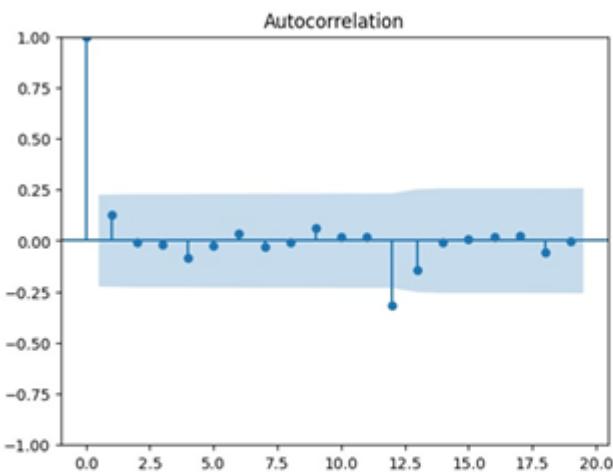
evaluate(test["Electricity_Consumption"], sarima_forecast, "SARIMA")
evaluate(test["Electricity_Consumption"], sarimax_forecast, "SARIMAX")
```

```
SARIMA Results
MAE : 5.56
RMSE: 7.17
-----
SARIMAX Results
MAE : 5.94
RMSE: 7.38
-----
```

```
In [13]: sarimax_result.resid.plot(title="SARIMAX Residuals")
plt.show()

plot_acf(sarimax_result.resid.dropna())
plt.show()
```





In []: